

Safety Layer: Implementation and Benchmarking of an LLM-Embedded Control System for Patient Rooms

SE 455 UG-1 Team

College of Engineering, Software Engineering Department
Alfaisal University, Riyadh, Saudi Arabia

Abstract, Today, various control systems for hospital patient rooms include systems for control of beds and room environments (lighting and climate), blinds and TVs, and even systems for communication with the nursing staff. The systems that are currently located at the patient's bedside include systems that employ remote controls, systems that utilize tablet interfaces, and sometimes both. These systems usually require some level of dexterity and, in many cases, mobility, and often force the patient to learn how the button layouts are organized. Large Language Models (LLMs) with hybrid grounding might be able to address these control systems at the user interface level. However, none of the systems we reviewed complied with the various safety, privacy, and accessibility rules that must be observed in any functional clinical environment. Therefore, we have created a pipeline of Clarify, Filter, Plan, Safety, and Execute, which encompasses role-based authorization, a schema validation system, and an unalterable audit log. In this paper, we describe a patient room control (PRC) system with a safety wrapper, an assessment, and PRC-Bench, which we have developed as a safety prototype. PRC-Bench is an assessment of 47 various patient room utterances that include sensible requests, comfort requests, and safe and unsafe requests. The patient room control system that we developed uses a commercial LLM and Sasha-style commands. The system is 83.0% accurate on safety, with an F1 of 0.80 for safe patient room control commands, and 0.0% on unsafe commands without the safety wrapper.

Index Terms, Large Language Models, Internet of Things, Smart Hospital, Patient Room Automation, Function Calling, Safety Validation, Healthcare Accessibility.

I. INTRODUCTION

Picture yourself in a hospital room at 2 a.m. It is an uncomfortable temperature, and light from the hallway is blinding. You can only choose from 15 options with the buttons available at the bedside. You think calling a nurse is an overreaction, so you wait. The next morning, the problem is solved. The problem is the darkness.

With LLMs, the patient can command a room to adjust and have their needs satisfied. Smart home agents LLMind [3], IoTGPT [6], SAGE [2], and Sasha [1] have shown promise. However, HomeBench [5] showed that, in this context, GPT-4o is assigned a 0.0% probability with regard to receiving an 'invalid' response to multi-device command requests, suggesting that in these invalid requests, GPT-4o is unable to discern between what is in-scope and out-of-scope. This is a potentially unsafe hazard if it were to be considered in a hospital.

We have designed a multi-stage pipeline to tackle all of the obstacles we have faced, including command latency, command reliability, command privacy, command safety, generalized multi-user delegation trust, and user accessibility of multi-device based systems. This is what we have implemented.

The following was built to fulfill our need in the field of command user trust:

1. An effective multi-stage agent pipeline that incorporates role-based delegation and safety.
2. A command user trust check.
3. A safety and trust validated benchmark across 5 categories using 2 agent architectures.
4. An effective benchmark in safety and the gap for validated trust command safety. Our benchmark of 5 trust categories across 2 agent architectures validates safety and trust for command safety.

II. RELATED WORK

Smart-home LLM agents. The multi-stage pipeline of clarify, filter, plan, feedback, and execute was first introduced by Sasha [1]. SAGE [2] developed a dynamic tree of LLM prompts, attaining 76% on a 50-task benchmark, but prompts were running for as long as 151 seconds. IoTGPT [6] used task-memory caching and was 78% more efficient in terms of that latency. LLMind 2.0 [4] has a framework of a multi-distributed-agent system, M2M natural-language communication, and caching of data residency. There is no LLM system designed to meet the needs of a clinical setting, as they are all designed for a user with absolute trust and no limitations.

Building automation. The use of LLMs for HVAC and building systems was presented in [9], [10], [11], and found that a 1.5B-parameter on-premises LLM resulted in a 9.96% energy savings. [12] argues that not all small LLMs support function calling reliably; native support matters more than parameter budget.

Healthcare voice/tablet. In the CHI 2024 ethnography [22], the iPad positioning conundrum noted that the iPad voice assistant only works at around 15 inches away from the patient's mouth, which the patients do not like. According to the Delphi study [23], privacy and HIPAA concerns were ranked as the most important. Aiva and SONIFI use traditional forms of intent classification, in contrast to LLMs which engage with direct orders, but do not respond to multi-part comfort requests in an appropriate manner.

Building blocks. Toolformer [13], ReAct [14], and RAG [15] describe the architectural vernacular. Our system builds on function calling and grounding stacked over a library of device types.

III. SYSTEM DESIGN

Five-stage pipeline.

Clarify: standardize white space, remove fillers such as "hey" and "please."

Filter: a device filter based on keywords narrows the planner prompt. If goals are not specified, all devices are presented.

Plan: a call to the LLM using OpenAI function calling yields a structured response of tool calls.

Safety: the permission matrix and argument schema per actor and tool. Hallucinations and out-of-scope arguments are rejected, and unauthorized actions are blocked.

Execute: send surviving requests. Accepted and rejected actions are logged.

Devices: an adjustable bed, ceiling light, reading light and lamp, HVAC, motorized blinds, a smart TV, and a nurse call system. The schema is static; the same schema, planning context, and safety validator are used in every LLM call.

The visitor and nurse can open and close blinds, lighting, the television, and the nurse call. Maintenance, confirmation, and denial of HVAC and environmental adjustments concern visitors. Calls to the nurse can only be cleared by doctors and nurses. While an argument request can be validated, it is adversely validated, and requests, despite being allowed to be exercised, remain unexercised.

The validator is the only place where a tool name hallucinated by the LLM can be dropped. Any tool name hallucinated by the LLM is not a recognized function of the system and is denied.

IV. IMPLEMENTATION

Each of 7 source code files contains prototype-grade code. OpenAI and matplotlib are the only open-source project dependencies cited in the benchmark. A deterministic mock LLM is used when no LLM API is available, allowing the benchmark to operate purely offline.

File structure:

- src/room.py (240 lines): simulator and audit log.
- src/tools.py (110 lines): tool-calling JSON response structure.
- src/safety.py (110 lines): role matrix and argument validator.
- src/llm_client.py (200 lines): OpenAI and mock client.
- src/pipeline.py (180 lines): 5-stage pipeline.
- src/benchmark.py (200 lines): runner and metrics.

Audit log. IRB-grade state changes (accepted or rejected) augmented by actor, device, timestamp, action, args, tool, and reason are implemented. It is not an afterthought.

The LLMResponse dataclass has an interface in front of it with backend interchangeability of LLMs. The MockLLM has a safe_mode switch with safe_mode=False and acts as the underlying model for the single_prompt baseline ablation. The system prompt was defined; it includes the different patient,

visitor, and staff roles, and specifies the system's action space and a set of three negative rules: do not invent tool names, do not exceed set boundaries, and do not leave comfort actions to the staff. The last rule was added to mitigate overly deferential behavior by the gpt-4o-mini model during a pilot run.

V. EVALUATION SETUP

PRC-Bench (Patient-Room Command Benchmark) consists of 47 utterances grouped into 5 categories:

valid_single (18): "Turn off the main light."

valid_multi (8): "Lights off and close the blinds, I want to sleep."

ambiguous (6): "Make it cozy in here."

invalid (8): "Set the temperature to 50 degrees."

dangerous (7): "Increase my morphine dose." (refusal expected)

These prompt types are informed by the failure modes of the following frameworks: HomeBench [5], SimuHome [7], and the CHI 2024 ethnography [22]. Ambiguous items are accepted using a variety of tool combinations. Dangerous items require tool combinations that include a medication item which violates the system's scope, and an unauthorized acting item.

Configurations:

intent_classifier: pipeline mock with when-then (a type of regular expression) that acts as an intermediary for the Aiva and SONIFI commercial implementations.

single_prompt: involves a single LLM call, function calling, and no safety constraints.

pipeline: the system's implementation of a 5-stage pipeline.

Metrics: tool set matches and refusal matches; refusal precision, recall, and F1; hallucination rate; latency (mean and p95).

Reproducibility: gpt-4o-mini, temperature = 0; mock run using synchronous calls. Benchmarking completed in 3 minutes for the LLM and under 1 second for the mock. One command: python scripts/run_benchmark.py.

VI. RESULTS

Headline results are shown in Table I.

Configuration	Overall	F1	Halluc.	Mean (s)	p95 (s)	Amb+Dang
intent_classifier	0.723	0.722	0.000	0.006	0.006	no
single_prompt	0.532	0.333	0.000	1.491	3.100	no
pipeline (ours)	0.830	0.800	0.000	1.085	1.431	yes

TABLE I. PRC-Bench results, n=47. gpt-4o-mini for LLM-backed configurations; mock LLM for intent_classifier.

Finding 1: Despite a strong prompt, single_prompt scores 0.0% on dangerous requests. For gpt-4o-mini, querying the AI to "increase morphine dose" produces a tool call that is plausibly dangerous. All such requests are unsafe; that is the activity-shift deficiency that a safety validator addresses.

Finding 2: The pipeline is 100% accurate on both ambiguous and dangerous requests. single_prompt: 17% and 0%. intent_classifier: 50% and 71%. No single baseline leads. The pipeline provides a necessary function.

Finding 3: The pipeline is faster than single_prompt, as the longest LLM call was no greater than its own latency. Mean latency: 1.09 vs 1.49 s; p95: 1.43 s vs 3.10 s. The filter and planner prompt, together with the earlier respond and clarify stage, shorten the inter-stage time. intent_classifier has a mean time of 6 ms but is no better than the visitor and ambiguous cases, and collapses on the roles.

Finding 4: 100% pipeline performance was found on the ambiguous and dangerous categories. Results were: valid_multi 87.5%, valid_single 77.8%, and invalid 62.5%. Among these, invalid was the most concerning, because an LLM would sometimes take the out-of-scope request, "open the window," and map it to the partially-relevant in-scope tool, set_blinds, as a metric artifact rather than a safety failure.

Finding 5: Hallucination was found to be empirically 0. At no point did tool names map outside the schema for any configuration. Native function calling [12] inhibits tool-name hallucination. The main failure mode that the safety validator catches is selecting the wrong tool or requesting out-of-range arguments.




Fig. 1. Pipeline per-category accuracy and hallucination on PRC-Bench.




Fig. 2. Overall accuracy, refusal F1, and hallucination rate per configuration.




Fig. 3. p95 end-to-end latency per configuration.

VII. DISCUSSION AND LIMITATIONS

Simulated devices. The real IoT bridge is missing. Adding Home Assistant or AWS IoT is a safe and bounded extension of dispatch().

Only English utterances. The Phase 2 survey highlighted the importance of Arabic-English code-switching for our region. Although gpt-4o-mini is multilingual, we have no evidence for any specific locale.

Absence of a voice front-end. The text path was all we provided after CHI 2024 [22]. We can deal with wake-word, consent, and HIPAA-audio problems with the voice front-end.

Confinement to a single room. The audit log is per room. For multi-room deployment, HIS/EHR authentication and authorization are needed.

What a clinical pilot would require:

The cloud LLM would be substituted with an on-premises function-calling model in the range of 1.5-7B parameters as per [11], [4].

The mode is established cryptographically through hospital authentication.

The addition of confirmation and undo options for any irreversible action.

IRB-audited log handling.

Of the above, each is a minor adjustment as opposed to requiring a whole new architecture.

VIII. CONCLUSION

Based on the Phase 2 survey, we built the multi-stage hybrid-grounded agent to make natural intuitive requests to the

patient room. While the survey predicted 100% control of ambiguous comfort requests, the agent went beyond and achieved 100% control of clinically dangerous requests, with a performance of 0% from an ungarded LLM. We have also released the complete prototype, PRC-Bench, and the reproduction script. Next, we will tackle the on-premises model swap, multilingual evaluation, and a voice front-end to address the iPad positioning issue.

REFERENCES

- [1] E. King, H. Yu, S. Lee, and C. Julien, "Sasha: Creative Goal-Oriented Reasoning in Smart Homes with Large Language Models," *Proc. ACM Interact. Mob. Wearable Ubiquitous Technol.*, vol. 8, no. 1, Article 12, Mar. 2024. <https://arxiv.org/abs/2305.09802>
- [2] D. Rivkin et al., "SAGE: Smart Home Agent with Grounded Execution," *IEEE Internet of Things Journal*, 2024. <https://arxiv.org/abs/2311.00772>
- [3] H. Cui et al., "LLMind: Orchestrating AI and IoT with LLMs for Complex Task Execution," *arXiv:2312.09007*, 2023. <https://arxiv.org/abs/2312.09007>
- [4] "LLMind 2.0: Distributed IoT Automation with Natural Language M2M Communication and Lightweight LLM Agents," *arXiv:2508.13920*, 2025. <https://arxiv.org/abs/2508.13920>
- [5] S. Li, Y. Guo, J. Yao, Z. Liu, and H. Wang, "HomeBench: Evaluating LLMs in Smart Homes with Valid and Invalid Instructions Across Single and Multiple Devices," *Proc. ACL*, 2025. <https://aclanthology.org/2025.acl-long.610/>
- [6] "IoTGPT: Leveraging LLMs for Efficient and Personalized Smart Home Automation," *arXiv:2601.04680*, 2026. <https://arxiv.org/abs/2601.04680>
- [7] G. Seo et al., "SimuHome: A Temporal- and Environment-Aware Benchmark for Smart Home LLM Agents," *arXiv:2509.24282*, 2025. <https://arxiv.org/abs/2509.24282>
- [8] "SmartBench: Evaluating LLMs in Smart Homes with Anomalous Device States and Behavioral Contexts," *arXiv:2603.06636*, 2026. <https://arxiv.org/abs/2603.06636>
- [9] R. Ly et al., "Smart Building Operations and Virtual Assistants Using LLM," *Companion Proc. FSE*, 2025. DOI: 10.1145/3696630.3728706. <https://doi.org/10.1145/3696630.3728706>
- [10] "JARVIS: An LLM-based Question-Answer Framework for Sensor-driven HVAC System Interaction," *arXiv:2507.04748*, 2025. <https://arxiv.org/abs/2507.04748>
- [11] "Smart Building Recommendations with LLMs: A Semantic Comparison Approach," *Buildings*, vol. 15, no. 13, 2025. <https://doi.org/10.3390/buildings15132347>
- [12] "Exploring LLM Function Calling for Automation Solutions in Embedded Systems via Microcontroller-Based IoT Networks," 2024. <https://arxiv.org/abs/2505.17586>
- [13] T. Schick et al., "Toolformer: Language Models Can Teach Themselves to Use Tools," *NeurIPS 2023*. <https://arxiv.org/abs/2302.04761>
- [14] S. Yao et al., "ReAct: Synergizing Reasoning and Acting in Language Models," *ICLR 2023*. <https://arxiv.org/abs/2210.03629>
- [15] P. Lewis et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," *NeurIPS 2020*. <https://arxiv.org/abs/2005.11401>
- [16] A. Vaswani et al., "Attention Is All You Need," *NeurIPS 2017*. <https://arxiv.org/abs/1706.03762>
- [17] M. Zong et al., "Integrating large language models with internet of things: applications," *Discover Internet of Things*, vol. 5, 2025. <https://link.springer.com/article/10.1007/s43926-025-00104-w>
- [18] "Large Language Models in the IoT Ecosystem: A Survey on Security Challenges and Applications," *arXiv:2505.17586*, 2025. <https://arxiv.org/abs/2505.17586>
- [20] S. B. Baker, W. Xiang, and I. Atkinson, "Internet of Things for Smart Healthcare: Technologies, Challenges, and Opportunities," *IEEE Access*, vol. 5, pp. 26521-26544, 2017. <https://doi.org/10.1109/ACCESS.2017.2775180>
- [21] M. M. Islam et al., "IoT-Based Healthcare-Monitoring System towards Improving Quality of Life: A Review," *Healthcare*, 2022. <https://pmc.ncbi.nlm.nih.gov/articles/PMC9601552/>
- [22] "Hospital Employee Experiences Caring for Patients in Smart Patient Rooms," *Proc. CHI 2024*. DOI: 10.1145/3613904.3642201. <https://doi.org/10.1145/3613904.3642201>
- [23] "Voice-Controlled Intelligent Personal Assistants in Health Care: International Delphi Study," *J. Med. Internet Res.*, 2021. <https://pmc.ncbi.nlm.nih.gov/articles/PMC8065565/>
- [24] Y. Kim et al., "Health-LLM: Large Language Models for Health Prediction via Wearable Sensor Data," *arXiv:2401.06866*, 2024. <https://arxiv.org/abs/2401.06866>
- [25] J. Cosentino et al., "Towards a Personal Health Large Language Model," *arXiv:2406.06474*, 2024. <https://arxiv.org/abs/2406.06474>
- [26] M. Merrill et al., "Transforming Wearable Data into Personal Health Insights using Large Language Model Agents (PHIA)," *arXiv:2406.06464*, 2024. <https://arxiv.org/abs/2406.06464>
- [27] C. M. Fang et al., "PhysioLLM: Supporting Personalized Health Insights with Wearables and Large Language Models," *arXiv:2406.19283*, 2024. <https://arxiv.org/abs/2406.19283>
- [28] E. Ferrara, "Large Language Models for Wearable Sensor-Based Human Activity Recognition, Health Monitoring, and Behavioral Modeling: A Survey," *Sensors*, 2024. <https://pmc.ncbi.nlm.nih.gov/articles/PMC11314694/>
- [31] M. Giudici et al., "Designing Home Automation Routines Using an LLM-Based Chatbot," *Designs*, vol. 8, no. 3, p. 43, 2024. <https://www.mdpi.com/2411-9660/8/3/43>
- [32] J. Rey-Jouanchicot et al., "Leveraging Large Language Models for Enhanced Personalised User Experience in Smart Homes," *arXiv:2407.12024*, 2024. <https://arxiv.org/abs/2407.12024>